



Optimization Techniques for Multi Cooperative Systems MCTR 1021
Mechatronics Engineering
Faculty of Engineering and Materials Science
German University in Cairo

TITLE

By

TeamNum

Member1

Member2

Member3

Member4

Member5

Course Team

SUPERVISOR1

SUPERVISOR2

SUPERVISOR3

December 11, 2025

This is to certify that:

- (i) the report comprises only my original work toward the course project,
- (ii) due acknowledgment has been made in the text to all other material used

TeamNum
December 11, 2025

Acknowledgments

I would like to thank the following for their efforts with me in my study and for tolerating me during this journey...

Abstract

The abstract of the document is added here...

It is usually written after finishing writing all the other chapters.

Contents

Acknowledgments	iii
List of Abbreviations	vii
List of Figures	viii
List of Tables	ix
1 Introduction	1
1.1 Section Name	1
1.2 Another Section	1
2 Methodology	9
2.1 Simulated Annealing Implementation	17
2.1.1 Pseudocode	18
2.1.2 Flowchart	19
2.1.3 Problem-Specific Considerations	19
2.1.4 GenerateNeighbor Subroutine	20
2.1.5 Key Algorithm Components	20
2.2 Genetic Algorithm Implementation	21
2.2.1 Pseudocode	24
2.2.2 Problem-Specific Considerations	25
2.3 Ant Colony Optimization Implementation	25
2.3.1 Pseudocode	26
2.3.2 Implementation Details	26
2.4 Artificial Bee Colony Implementation	27
2.4.1 Algorithm Flow	27
2.4.2 Key Operators and Parameters	28
2.4.3 Pseudocode	29
3 Results	30

3.1	Simulated Annealing Results	30
3.1.1	Parameters Tuning	30
3.1.2	Cooperative Robot Number Configurations	32
3.1.3	Visualization	32
3.1.4	Conclusions	34
3.2	Genetic Algorithm Results	35
3.2.1	Parameters Tuning	35
3.2.2	Cooperative Robot Number Configurations	37
3.2.3	Testing Crossover, Mutation and Parent Selection Methods	39
3.2.4	Visualization	42
3.2.5	Conclusions	44
3.3	Ant Colony Optimization Results	44
3.3.1	Parameters Tuning	44
3.3.2	Case Study Comparisons (SA vs GA vs ACO vs ABC)	45
3.3.3	Performance Indices	48
3.4	Algorithm Validation	49
3.4.1	Time Step Equivalence Calculation	49
3.4.2	Results Comparison	50
4	Conclusion	51
5	Future Work	52
	Appendix	53
	References	54

List of Abbreviations

MRSs	Multi-Robot Systems
UAVs	Unmanned Aerial Vehicles
UGVs	Unmanned Ground Vehicles

List of Figures

1.1	GUC Logo	1
2.1	High-level overview of the path planning optimization approach.	9
2.2	Simulated Annealing algorithm flowchart for multi-robot path planning.	19
3.1	Graphical visualization of the optimization process, showing robot paths and the area covered in the environment.	33
3.2	Terminal visualization and coverage results, displaying key performance metrics.	34
3.3	Graphical visualization of the optimization process for the best case study C1, showing robot paths and the area covered in the environment.	43
3.4	Terminal visualization and coverage results, displaying key performance metrics.	43
3.5	Convergence comparison for Case 1 (6 Robots, 125 Steps). ACO starts with a superior solution due to heuristic initialization and maintains the lead throughout.	46
3.6	Convergence comparison for Case 3 (3 Robots, 50 Steps). ACO demonstrates rapid convergence from a high initial cost to outperform the other algorithms within the first 10 iterations.	46
3.7	Convergence comparison for Case 4 (4 Robots, 100 Steps, 300 Iterations). The extended iterations show ACO maintaining its lead and stability.	47
3.8	Convergence comparison for Case 10 (2 Robots, 100 Steps). In this challenging scenario with minimal agents, ACO's advantage is particularly pronounced from the start.	47

List of Tables

3.1	Performance Metrics: Area Coverage Results Summary	30
3.2	Performance Metrics for Different Robot Counts	32
3.3	Performance Metrics: Area Coverage Results Summary for mutation rate = 30% and elite rate 10% while varying Population and Generation Size	35
3.4	Performance Metrics: Area Coverage Results Summary for mutation rate = 30% and elite rate 10% while varying α and β	35
3.5	Performance Metrics: Area Coverage Results Summary for Population size = 20 and Generation Size 200.	36
3.6	Performance Metrics for Different Robot Counts	37
3.7	Here we use all the parameters same as C1 but vary the Mutation, Crossover and Parent Selection Methods.	39
3.8	Genetic Operator Methods Summary	40
3.9	ACO Performance Metrics (AS vs SACO Variants)	44
3.10	Case Study Results Across SA, GA, ACO, and ABC	45
3.11	Function Minimization: Performance Indices (20 Runs)	48
3.12	Comparison of Coverage Performance with Benchmark Study	50

Chapter 1

Introduction

1.1 Section Name

Some sample text with an Unmanned Aerial Vehicles (UAVs), some citation [?, ?], and some more Unmanned Ground Vehicles (UGVs).

1.2 Another Section

Reference to Section 1.1, and reuse of UGVs nad UAVs with also full use of Multi-Robot Systems (MRSs). Reference to figure 1.1.



Figure 1.1: GUC Logo

Shen *et al.* (2025) [1] introduced **Multi-CAP**, a hierarchical algorithm for multi-robot coverage path planning in unknown environments. The main objective of their approach is to enable multiple robots to cooperatively explore and cover a given area while minimizing redundant paths and maintaining network connectivity among robots.

The **problem is formulated** as an adjacency graph $G = (N, E)$, where N represents the subareas to be covered and E denotes the possible travel paths between them. Each robot starts from a known position and is assigned a sequence of subareas to visit, ensuring that all regions are covered with minimal total travel distance.

The **decision variables** define which subareas are assigned to each robot and in what visiting order. Each route must begin at the robot's starting point and terminate at its final visited subarea, without requiring a return to the start (open route).

The **objective function** aims to minimize the total travel cost of all robots:

$$J(\Pi) = \sum_{m=1}^M \sum_{k=0}^{K_m-1} w_{n_m(k), n_m(k+1)} \quad (1.1)$$

where M is the number of robots, K_m is the number of subareas assigned to robot m , and w_{ij} represents the travel cost (distance) between subareas i and j .

The problem is subject to several **constraints**: (1) every subarea must be visited exactly once (coverage constraint), (2) no subarea is assigned to more than one robot (assignment constraint), (3) all robot movements must follow valid edges in E (connectivity constraint), and (4) each robot's path must follow an open-route structure.

Shen *et al.* applied a two-stage heuristic **optimization approach**. In the first stage, a *Nearest-Neighbor heuristic* assigns subareas to robots based on proximity, generating initial feasible routes with minimal overlap. In the second stage, a *2-Opt local search algorithm* refines each route by iteratively adjusting the visiting order to reduce total travel distance.

Finally, each robot performs local coverage within its assigned region using a back and forth sweeping motion or by solving a small local Traveling Salesman Problem (TSP) with the 2-Opt method to ensure complete coverage. This hierarchical framework achieves efficient, connectivity aware exploration while minimizing total path length and coverage time.

In a different study, Mikkelsen *et al.* (2024) [2] presented a novel approach for solving the Multi-Robot Communication-Aware Trajectory Planning (MR-CaTP) problem, which reformulates a general optimization framework by using changes in robot positions as decision variables and introducing linear constraints to ensure communication performance via the Fiedler value (the second smallest eigenvalue of the graph Laplacian) and collision avoidance.

The **problem is formulated** as a receding horizon optimization over K steps for a swarm

of N holonomic robots in $\mathbb{R}^n$, where the communication network is modeled as a time-varying undirected graph $G(t) = (\mathcal{V}, E(t))$ with edge weights $w_{ij}(t)$ representing link quality based on inter-robot distances. The goal is to compute feasible and optimal trajectories that maintain network connectivity while fulfilling task-specific objectives, such as inspection or arbitrary motion with communication insurance.

The **decision variables** are the concatenated sequence of position changes (inputs) over K future steps, $\mathbf{U}^k = (\mathbf{u}^k, \mathbf{u}^{k+1}, \dots, \mathbf{u}^{k+K-1}) \in \mathbb{R}^{KnN}$, where each $\mathbf{u}_i^{k+h}$ defines the change in position for robot i at step $k+h$.

The **objective function** minimizes the sum of a generic cost over the K steps:

$$\min_{\mathbf{U}^k} \sum_{h=0}^{K-1} f(\mathbf{u}^{k+h}), \quad (1.2)$$

where $f(\cdot)$ is task-dependent (e.g., quadratic error to points of interest in inspection tasks or deviation from reference inputs in communication insurance).

The problem is subject to several **constraints**: (1) predicted kinematics: $\mathbf{P}^{k+1} = \mathbf{1}_{K \times 1} \otimes \mathbf{p}^k + \mathbf{B}\mathbf{U}^k$, where $\mathbf{P}^{k+1}$ accumulates future positions and $\mathbf{B}$ is a lower triangular accumulation matrix, (2) input bounds: $\|\mathbf{u}_i^{k+h}\|_p \leq \bar{u}$ for all robots i and steps h , ensuring bounded speed, (3) approximated communication performance: $\mathbf{M}^k \mathbf{U}^k \geq (\underline{\lambda}_2 - \lambda_2^k) \mathbf{1}_{K \times 1}$, a linear perturbation-based constraint to keep predicted Fiedler values $\hat{\Lambda}_2^{k+1}$ above a lower bound $\underline{\lambda}_2 > 0$ for connectivity, and (4) approximated collision avoidance: $\tilde{\mathbf{C}}^k \mathbf{U}^k \leq \tilde{\mathbf{d}}^k$, a linear inequality based on buffered Voronoi partitions from the Delaunay triangulation to prevent inter-robot overlaps.

Mikkelsen *et al.* applied this framework in a **general optimization approximation** using first-order Taylor expansion for Fiedler value prediction and Voronoi-based linearization for tractability, enabling significant computational speedup while remaining close to optimal solutions. In the inspection task application, points of interest are first assigned via minimum perfect matching (solved with the Hungarian algorithm), then trajectories are optimized with a quadratic objective incorporating error to assigned points and Fiedler gradient for relays. In the communication insurance service, the objective minimizes deviation from desired reference inputs, with optional slack variables to soften the Fiedler bound when strict feasibility is challenging. Simulations validate that the Fiedler constraint ensures connectivity at all times, with approximations reducing solve times effectively.

Furthermore, Kashyap *et al.* (2023) [3] introduced a hybrid Lévy Flight–Particle Swarm Optimization (LF-PSO) algorithm for multi-robot exploration in Urban Search and Rescue (USAR) scenarios. Unlike previous approaches that assume known maps, global positioning, or unrestricted communication, the LF-PSO model operates in unknown, GPS-denied, and communication-limited environments. The algorithm combines Lévy Flight for random, long-range exploration with PSO-based inter-robot repulsion to prevent clustering and improve

coverage. Simulation results in a 20×20 m area showed that LF-PSO achieved about 76% area coverage, outperforming standalone PSO and LF methods. The study demonstrates a robust, decentralized approach for autonomous robot coordination in real-world search and rescue tasks.

The **problem is formulated** under assumptions that no initial information about the environment (size or obstacle locations) is known, no global positioning or information about the location of other robots is available, no information about the target or victim's position is known, and communication is only possible when other robots are within a specified communication radius d . Initialization involves randomizing the position and velocity (within the constraints) of the robots in the search space and setting the number of iterations where each iteration represents a time step (e.g., robot at time step $t+1$ has moved when comparing with respect to robot at time step t). To facilitate comprehensive area coverage, each robot's path is represented as a sequence of waypoints in the two-dimensional search space. Waypoints are interconnected to form a path that guides the robot through the search area. The order in which waypoints are connected determines the robot's trajectory. Each waypoint maintains an exploration status to indicate whether it has been visited by the robot or remains unexplored.

The **decision variables** define the best position of waypoints to achieve most efficient exploration. The main decision variables are derived from the hybrid Lévy Flight - Particle Swarm Optimization (LF-PSO) update equations: $x_i(t), y_i(t)$: Position of robot i at time t . $V_{i,j}(t)$: Velocity of robot i in direction j at time t . s_i : Step size in Lévy Flight movement. $L\theta_i$: Movement direction in Lévy Flight. $R\theta_i$: Repulsion direction from nearby robots. P_{lbx_i}, P_{lby_i} : Local best positions (from Lévy Flight exploration), where local means the decision is only based on the Levi Flight Step Distance, not taking any cooperative information from other robots, which affects their motion such as the repulsion mechanism. P_{gbx_i}, P_{gby_i} : Global best positions (from repulsive PSO component).

The **objective function** aims to maximize the total exploration efficiency (fraction of the area explored by all robots):

$$\text{Maximize } f = \frac{A_{\text{covered}}}{A_{\text{total}}} \quad (1.3)$$

where A_{covered} = Total area explored by all robots collectively, and A_{total} = Total area of the environment.

The problem is subject to several **constraints**: (1) Robots must remain within the map boundaries at all times, (2) Each robot's velocity V_i and step size s_i are bounded by maximum limits: $0 < V_i \leq V_{\text{max}}, 0 < s_i \leq s_{\text{max}}$, (3) Step Size Limit (Based on the Levy Flight algorithm), where the Lévy flight algorithm is a type of random walk used in optimization and search algorithms that combines frequent short-distance steps with occasional long-distance jumps (also randomizes motion direction L_θ from), (4) Communication between robots occurs only when they are within a communication range d , (5) Obstacles are detected only when

within a certain proximity based on the robot's front, left, and right sensors, (6) Movement in the presence of obstacles follows the policy: Turn $\frac{\pi}{2}$ Left/Right if Left Sensor FALSE, Front Sensor TRUE, Right Sensor FALSE; Turn $\frac{\pi}{2}$ Left if Left Sensor FALSE, Front Sensor TRUE, Right Sensor TRUE; Turn $\frac{\pi}{2}$ Right if Left Sensor TRUE, Front Sensor TRUE, Right Sensor FALSE; Turn π if Left Sensor TRUE, Front Sensor TRUE, Right Sensor TRUE; Turn $\frac{\pi}{2}$ Right if Left Sensor TRUE, Front Sensor FALSE, Right Sensor FALSE; Turn $\frac{\pi}{2}$ Left if Left Sensor FALSE, Front Sensor FALSE, Right Sensor TRUE, (7) The repulsion distance D_j between robots 1 and 2 is given by $D_j = \frac{1}{d_1} + \frac{1}{d_2}$, (8) Do not pass by an explored way-point. In mathematical formulation, the robot's motion is governed by: $P_{lb_{x_j}} - x_j = s_j \cos(L\theta_j)$, $P_{lb_{y_j}} - y_j = s_j \sin(L\theta_j)$, $P_{gb_{x_j}} - x_j = D_j \cos(R\theta_j)$, $P_{gb_{y_j}} - y_j = D_j \sin(R\theta_j)$. The modified PSO velocity and position update equations are (this should change when we use other optimization algorithms): $V_{i,j}(t+1) = \omega V_{i,j}(t) + p_\omega \text{rand}(P_{lb_j}(t) - x_{i,j}(t)) + n_\omega \text{rand}(P_{gb_j}(t) - x_{i,j}(t))$, $x_{i,j}(t+1) = x_{i,j}(t) + V_{i,j}(t+1)$, where ω is the inertia weight, p_ω and n_ω are cognitive and social coefficients respectively.

Kashyap *et al.* applied a **hybrid LF-PSO optimization approach** that integrates Lévy Flight for exploration with PSO for repulsion-based coordination, updating robot positions and velocities iteratively in a decentralized manner to maximize coverage while adhering to constraints in unknown environments.

In addition, Som *et al.* (2024) [4] proposed a Particle Swarm Optimization (PSO) based approach for relay-node placement in cluster-based wireless sensor networks (WSNs). The problem was formulated as a multi-objective optimization problem that maximizes connectivity between the placed relay nodes (RNs), maximizes the coverage of sensor nodes (SNs) by the placed relay nodes, and minimizes the number of relay nodes deployed. The main challenge was to strategically position RNs from a set of candidate locations to act as cluster heads. This optimized placement approach balances both cost-effectiveness and performance requirements to enable efficient network backbone construction. This relay node infrastructure deployment exemplifies how *meta-heuristic optimization techniques* can address complex coordination challenges in networked systems, where distributed agents must collectively satisfy multiple objectives while maintaining communication constraints—a fundamental problem that parallels the multi-robot path planning and trajectory coordination strategies.

The **problem is formulated** as an application environment where N sensor nodes $S = \{s_1, s_2, \dots, s_n\}$ are deployed randomly or manually across a monitored region, and M potential positions $P = \{p_1, p_2, \dots, p_m\}$ are available for relay node placement. The task is to determine which candidate positions should receive relay nodes to form a cluster-based network.

The **decision variables** include : δ_i indicating whether position p_i is selected for relay node deployment, ω_i indicating whether sensor node s_i is covered by at least one relay node within communication range, ϕ_{ij} indicating connectivity between relay nodes r_j and r_i within communication distance, and α_i indicating whether relay node r_i can directly reach the base station.

There are three **objective functions**. The **first objective function** minimizes

$$F_1 = \frac{1}{M} \sum_{i=1}^M x_i \quad (1.4)$$

to reduce the total number of deployed relay nodes and associated costs. The **second objective function** maximizes

$$F_2 = \sum_{\forall s_i \in S} Cov(s_i) \quad (1.5)$$

to ensure comprehensive sensor coverage across the monitored region, where $Cov(s_i)$ represents the set of relay nodes within communication range of sensor node s_i . The **third objective function** maximizes

$$F_3 = \sum_{\forall r_i \in R} Con(r_i) \quad (1.6)$$

to establish robust connectivity among relay nodes forming a reliable backbone toward the base station, where $Con(r_i)$ denotes the set of relay nodes within communication range of r_i and positioned toward the BS. These three objectives are integrated into a composite **fitness function**:

$$\text{Maximize } Fitness = w_1(1 - F_1) + w_2F_2 + w_3F_3 \quad (1.7)$$

where the weights satisfy $\sum w_i = 1$ and $0 \leq w_i \leq 1$ for $i = 1, 2, 3$.

The problem is subject to several **constraints** : (1) coverage constraint ($\omega_n = 1, \forall n, 1 \leq n \leq N$), which ensures that each sensor node has minimum one relay node within its communication range S_{com} , thereby enabling each sensor to be part of a cluster, (2) connectivity constraint ($\sum_{j=1}^Z (\phi_{ij} \times \delta_j + \alpha_i) \times \delta_i \geq 1, \forall i, 1 \leq i \leq Z$), which ensures that all deployed relay nodes maintain connectivity with the base station, (3) binary constraint for relay placement ($\delta_i \in \{0, 1\}, \forall i, 1 \leq i \leq M$), which restricts relay node selection to binary values, (4) binary constraints for connectivity and base station reach ($\alpha_i \in \{0, 1\}, \forall i, 1 \leq i \leq Z$ and $\phi_{ij} \in \{0, 1\}, i \neq j, \forall i, j, 1 \leq i \leq Z, 1 \leq j \leq Z$), which ensure that connectivity relationships are strictly defined.

Som *et al.* applied **Particle Swarm Optimization (PSO)** to solve this constrained multi-objective problem, where each particle encodes a candidate relay placement configuration. The PSO algorithm iteratively updates particle velocities and positions guided by personal and global best solutions, refining placements through a balance of exploration and exploitation. Simulation results across two network scenarios demonstrated that PSO selected fewer relay nodes than the Gravitational Search Algorithm while maintaining complete coverage and connectivity, validating its effectiveness for communication-aware relay node deployment.

Finally, Huang *et al.* (2020) [5] addressed a critical gap in multi-robot coverage path planning (MCP) research by proposing an algorithm specifically designed for outdoor environ-

ments with multiple land cover types. This work builds upon and extends existing MCPP methodologies to account for environmental variation, that was overlooked in prior research.

The **problem is formulated** as maximizing the coverage of a rectangular area D representing the search and rescue region, and s robots $\mathcal{R} = \{R_1, R_2, \dots, R_s\}$ with initial positions $\mathcal{S}_p = \{S_{p_1}, S_{p_2}, \dots, S_{p_s}\}$. The environment contains multiple land cover types C_v (e.g., farm, road, grass, wood, bare land, Gobi). Each land cover affect the robot performance:

- Visual field $Fd_{C_{v_i}}$: the detection range perpendicular to movement direction
- Moving speed $Sd_{i,r,c}$: the velocity when traversing cell $C_{i,r,c}$

To accommodate these varying capabilities, we decompose the task area using a hierarchical quadtree structure with l layers, where:

$$l = 1 + \log_2 \left(\frac{Fd_{\max}}{Fd_{\min}} \right) \quad (1.8)$$

The quadtree construction follows a two-phase process:

Phase 1: Top-down decomposition. Starting from the top layer with cell size $Sc_1 = 2 \times Fd_{\max}$, each subsequent layer i has cells of size $Sc_i = Sc_{i-1}/2$, creating cells $C_{i,r,c}$ at all layers.

Phase 2: Bottom-up integration. To eliminate redundant coverage, cells are classified as either *active* (assigned a land cover type) or *auxiliary* (set to null). Starting from the bottom layer l :

- Each cell $C_{l,r,c}$ is assigned the dominant land cover type within its boundaries
- For layers $i = l - 1$ to 1: If four child cells share the same non-null cover type AND their size is smaller than the robot's visual field in that terrain, they are integrated upward by:

$$C_{v_{i,r,c}} = C_{v_{i+1,r*2,c*2}} \text{ (parent inherits type)} \quad (1.9)$$

$$C_{v_{i+1,r*2,c*2}} = C_{v_{i+1,r*2+1,c*2}} = \dots = \text{null (children deactivated)} \quad (1.10)$$

This ensures all active cells $\mathcal{C}_{\text{active}} = \{C_{i,r,c} : C_{v_{i,r,c}} \neq \text{null}\}$ cover D exactly once, adapting cell sizes to local terrain complexity. To enable efficient task assignment, we establish adjacency relationships using shared neighbor directions $Dn_{i,r,c} \in \{0, 1, \dots, 15\}$ encoded in binary, where each bit represents a cardinal direction. For cells in different layers, shared neighbors are identified through logical AND operations on direction values.

The **decision variables** include :

- A_j : Area assigned to robot R_j
- P_j : Coverage path for robot R_j

The **objective function** minimizes the maximum task cost across all robots:

$$\min \max_{j \in [1, s]} \{T_j\} \quad (1.11)$$

where the task cost Qa_j for robot R_j is:

$$Qa_j = \sum_{C_{i,r,c} \in A_j} \frac{2Sc_i}{Sd_{i,r,c}} \quad (1.12)$$

The problem is subject to several **constraints**:

$$\bigcup_{j=1}^s A_j = D \quad (\text{Complete coverage}) \quad (1.13)$$

$$A_j \cap A_k = \emptyset, \quad \forall j, k \in [1, s], j \neq k \quad (\text{No overlap}) \quad (1.14)$$

$$|T_1| \approx |T_2| \approx \dots \approx |T_s| \quad (\text{Balanced workload}) \quad (1.15)$$

$$St_j \text{ is connected}, \quad \forall j \in [1, s] \quad (\text{Continuous}) \quad (1.16)$$

$$S_{p_j} \in A_j, \quad \forall j \in [1, s] \quad (\text{Initial position}) \quad (1.17)$$

Chapter 2

Methodology

Problem Formulation




Figure 2.1: High-level overview of the path planning optimization approach.

This chapter presents the mathematical formulation of the multi-robot exploration problem with distance-weighted connectivity constraints and energy limitations. We model the environment as a discrete grid and formulate the problem as a constrained optimization that balances area coverage with maintaining communication links between robots, while accounting for battery

constraints and penalizing undesirable behaviors such as network disconnection and obstacle encounters.

Assumptions

The environment is represented as a grid map $M \in \{0, 1, 2, 3\}^{n \times m}$ with unit cells of size 1×1 . Each cell $M[x, y]$ takes values: $M[x, y] = 0$ means unexplored, $M[x, y] = 1$ means free space, $M[x, y] = 2$ means obstacle, and $M[x, y] = 3$ marks a cell occupied by a robot.

We assume that all robots are exactly identical in their capabilities, including sensing range, communication radius, motion constraints, and battery capacity. This makes coordination simpler and ensures fairness in task allocation, as no robot has advantages over others.

We assume that obstacles are unknown at the start, which matches realistic exploration scenarios where the environment structure is discovered gradually. As robots explore, newly detected obstacles are shared among all robots through their communication network. Let $O_t \subseteq \{(x, y) \mid 1 \leq x \leq n, 1 \leq y \leq m\}$ denote the set of known obstacle locations at time t , with $O_t \subseteq O_{t+1}$ for all $t \geq 0$. This monotonic growth of obstacle knowledge captures the fact that once an obstacle is discovered, it remains known. When a robot encounters an obstacle during execution, it uses a reactive obstacle avoidance strategy to navigate around the obstruction while trying to maintain its overall trajectory toward the planned goal.

Robots are modeled as point agents that move synchronously at discrete time steps $t \in \{0, 1, 2, \dots\}$. Motion follows a 4-connected grid topology (up, down, left, right) or allows robots to stay in place. This discretization simplifies planning while maintaining sufficient capability for typical robotic platforms.

Each robot is equipped with a battery that permits movement for a maximum of $\ell \in \mathbb{N}$ cells, where $\ell > 0$ represents the energy budget. Since all robots are identical, they share the same energy capacity. This constraint reflects the practical limitation that robots must complete their exploration mission within their available energy resources.

We assume robots share their positions at each time step through local communication. The desirability of maintaining communication between any pair of robots is modeled by a weight function that decreases with their Euclidean distance, reflecting the fact that closer robots have more reliable communication links.

Map, Robots, and Positions

There are $N \in \mathbb{N}$ robots indexed by $i \in \{1, \dots, N\}$. The position of robot i at time t is denoted

$$p_{i,t} = (x_{i,t}, y_{i,t}) \in \mathbb{Z}^2,$$

where

$$x_{i,t} \in \{1, \dots, n\}, \quad y_{i,t} \in \{1, \dots, m\}.$$

Initial positions $p_{i,0}$ are given for all $i \in \{1, \dots, N\}$.

The discrete grid representation allows us to formulate the problem as a mixed-integer program. Each robot's state at any time is fully described by its integer coordinates within the map boundaries.

The 4-connected motion constraint is enforced by the Manhattan distance:

$$|x_{i,t} - x_{i,t-1}| + |y_{i,t} - y_{i,t-1}| \leq 1, \quad \forall i \in \{1, \dots, N\}, \forall t \geq 1.$$

This constraint ensures that each robot can move at most one grid cell per time step in one of the four cardinal directions, or remain stationary. The Manhattan distance metric naturally captures the 4-connected neighborhood structure and prevents diagonal movements, which simplifies motion planning and collision avoidance.

Energy Consumption Model

To model battery constraints, we define the cumulative distance traveled by robot i up to time t as:

$$D_{i,t} = \sum_{\tau=1}^t (|x_{i,\tau} - x_{i,\tau-1}| + |y_{i,\tau} - y_{i,\tau-1}|), \quad \forall i \in \{1, \dots, N\}, \forall t \geq 1.$$

Each unit of movement (one grid cell) consumes one unit of battery. The energy constraint for each robot is:

$$D_{i,t} \leq \ell, \quad \forall i \in \{1, \dots, N\}, \forall t \geq 0,$$

where $\ell \in \mathbb{N}$ is the maximum number of cells each robot can traverse.

This formulation ensures that the total distance traveled by any robot never exceeds its battery capacity. The constraint applies at every time step, meaning that robots must stop moving once they have exhausted their energy budget. Since robots are identical, they all share the same energy limit ℓ . This constraint forces the planner to balance exploration goals with energy availability, potentially requiring some robots to remain stationary or return to already-explored areas to conserve energy for critical coordination tasks. The mission effectively ends when all robots have either exhausted their batteries or completed the exploration task.

Coverage and Distance-Weighted Connectivity

Coverage is a fundamental metric in exploration tasks, measuring the fraction of the environment that has been visited by at least one robot. We define the binary coverage variable

$C_{x,y} \in \{0, 1\}$ for each cell (x, y) :

$$C_{x,y} = \begin{cases} 1, & \text{if } \exists i \in \{1, \dots, N\}, \exists t \geq 0 \text{ such that } p_{i,t} = (x, y), \\ 0, & \text{otherwise.} \end{cases}$$

The coverage is enforced via indicator constraints:

$$C_{x,y} \geq \mathbb{I}[p_{i,t} = (x, y)], \quad \forall x \in \{1, \dots, n\}, \forall y \in \{1, \dots, m\}, \forall i \in \{1, \dots, N\}, \forall t \geq 0,$$

where $\mathbb{I}[\cdot]$ is the indicator function that equals 1 when its argument is true and 0 otherwise.

These constraints ensure that once a cell is visited by any robot at any time, its coverage variable is set to 1 and remains 1 thereafter. This formulation naturally handles the cumulative nature of coverage: a cell needs to be visited only once to be considered covered, regardless of whether robots later leave or return to it.

To model communication quality between robots, we first define the Euclidean distance between robots i and j at time t as:

$$d_{ij,t} = \sqrt{(x_{i,t} - x_{j,t})^2 + (y_{i,t} - y_{j,t})^2} \in \mathbb{R}_{\geq 0}.$$

Building on this, we define the distance-weighted link quality $W_{ij,t} \in [0, 1]$ as:

$$W_{ij,t} = \begin{cases} \frac{1}{1 + d_{ij,t}}, & \text{if } d_{ij,t} \leq R_c, \\ 0, & \text{if } d_{ij,t} > R_c, \end{cases}$$

where $R_c > 0$ is the communication radius threshold.

This formulation captures two important aspects of real-world communication: (1) link quality degrades continuously with distance within communication range, following an inverse relationship, and (2) beyond a critical distance R_c , communication becomes impossible. The weight $W_{ij,t}$ reaches its maximum value of 1 when robots are at the same location ($d_{ij,t} = 0$) and decreases smoothly as distance increases, encouraging the optimizer to keep robots reasonably close while exploring.

Network Connectivity and Disconnection Penalty

Define the edge set at time t as:

$$E_t = \{(i, j) \mid i, j \in \{1, \dots, N\}, i < j, d_{ij,t} \leq d_{\text{threshold}}\}, \quad \forall t \geq 0,$$

where $d_{\text{threshold}} > 0$ is the connectivity distance threshold.

Let $G_t = (V, E_t)$ be the communication graph at time t , where $V = \{1, \dots, N\}$ is the set of robot indices. We define $\mathcal{C}_t$ as the set of connected components in G_t . For each robot i , let $C_t(i)$ denote the connected component containing robot i at time t .

To measure network disconnection, we introduce a binary indicator variable:

$$\delta_{i,t} = \begin{cases} 1, & \text{if robot } i \text{ is disconnected from the main network at time } t, \\ 0, & \text{otherwise,} \end{cases}$$

where the "main network" is defined as the largest connected component in G_t (breaking ties arbitrarily). Formally:

$$\delta_{i,t} = \mathbb{I} \left[C_t(i) \neq \arg \max_{C \in \mathcal{C}_t} |C| \right], \quad \forall i \in \{1, \dots, N\}, \forall t \geq 0.$$

When a robot is disconnected, we measure its distance to the main network. Let $V_{\text{main},t} \subseteq V$ denote the set of robots in the main network at time t . The distance from disconnected robot i to the main network is:

$$d_{i,\text{net},t} = \min_{j \in V_{\text{main},t}} d_{ij,t}, \quad \forall i \in \{1, \dots, N\} \text{ with } \delta_{i,t} = 1.$$

Since the mission continues until all robots exhaust their energy or complete exploration, we compute the total disconnection penalty by summing over all robots and all time steps up to the maximum mission duration T_{max} :

$$P_{\text{disconnect}} = \sum_{t=0}^{T_{\text{max}}} \sum_{i=1}^N \delta_{i,t} \cdot d_{i,\text{net},t},$$

where $T_{\text{max}} = \max\{\ell, nm\}$ is an upper bound on mission duration (either all energy is depleted or all cells are explored).

This penalty term has two important properties: (1) it is zero when all robots remain connected throughout the mission, and (2) when disconnection occurs, the penalty scales linearly with how far the isolated robot(s) are from the main network. This encourages the planner to either maintain full connectivity or, if temporary disconnection is unavoidable, to minimize both the duration and severity (distance) of the disconnection. Robots that are only slightly beyond communication range incur smaller penalties than those that have wandered far from the team.

Obstacle Encounter Penalty

To discourage robots from planning paths that lead to obstacle encounters, we define an obstacle encounter indicator:

$$\eta_{i,t} = \begin{cases} 1, & \text{if robot } i \text{ encounters an obstacle at time } t, \\ 0, & \text{otherwise.} \end{cases}$$

An obstacle encounter occurs when a robot plans to move to a cell that is later discovered to contain an obstacle. Formally, robot i encounters an obstacle at time t if:

$$p_{i,t} \in O_t \setminus O_{t-1}, \quad \text{for } t \geq 1.$$

This captures newly discovered obstacles that were unknown during planning. When such an encounter occurs, the robot must reactively navigate around the obstacle, leading to suboptimal paths, wasted energy, and delays.

The total obstacle encounter penalty is:

$$P_{\text{obstacle}} = \sum_{t=1}^{T_{\max}} \sum_{i=1}^N \eta_{i,t}.$$

This penalty counts the total number of obstacle encounters across all robots and all time steps. While it is impossible to completely avoid obstacles that are initially unknown, this penalty term encourages the planner to favor paths through regions with lower perceived obstacle density (based on prior exploration) and to coordinate robot movements such that scouts discover obstacles ahead of the main group. The penalty also indirectly encourages more conservative exploration strategies that reduce the likelihood of unexpected collisions.

Objective Function

The optimization problem seeks to maximize a composite objective that balances exploration, connectivity, energy efficiency, and safety:

$$\max f = \underbrace{\alpha \sum_{x=1}^n \sum_{y=1}^m C_{x,y}}_{\text{coverage term}} + \underbrace{\beta \sum_{t=0}^{T_{\max}} \sum_{i=1}^{N-1} \sum_{j=i+1}^N W_{ij,t}}_{\text{distance-weighted connectivity term}} - \underbrace{\gamma \cdot P_{\text{disconnect}}}_{\text{disconnection penalty}} - \underbrace{\zeta \cdot P_{\text{obstacle}}}_{\text{obstacle encounter penalty}},$$

where $\alpha, \beta, \gamma, \zeta \in \mathbb{R}_{>0}$ are positive weighting parameters.

The first term rewards exploring new areas by counting the total number of covered cells. The

second term encourages robots to maintain proximity throughout the mission by accumulating link quality weights over all time steps and all robot pairs. The third term penalizes network fragmentation, with penalties scaled by the distance of disconnected robots from the main network. The fourth term discourages obstacle encounters, which lead to inefficient reactive behaviors.

The parameters $\alpha, \beta, \gamma, \zeta$ allow us to tune the relative importance of these competing objectives. Larger α prioritizes aggressive exploration potentially at the cost of team fragmentation and safety. Larger β keeps the team cohesive but may slow down exploration. Larger γ enforces stricter connectivity requirements, particularly penalizing scenarios where robots become isolated far from the team. Larger ζ promotes more cautious exploration strategies that reduce collision risk. The choice of these parameters depends on mission requirements: search-and-rescue operations might prioritize coverage (α) and connectivity (β, γ), while mapping missions in hazardous environments might additionally emphasize safety (ζ).

Constraints

The optimization is subject to several constraints that ensure feasible and safe robot trajectories.

1. Map bounds:

$$x_{i,t} \in \{1, \dots, n\}, \quad y_{i,t} \in \{1, \dots, m\}, \quad \forall i \in \{1, \dots, N\}, \forall t \geq 0.$$

These constraints ensure that all robots remain within the boundaries of the defined map at all times. Violating these bounds would place robots in undefined regions outside the environment.

2. Motion constraint (4-connectivity):

$$|x_{i,t} - x_{i,t-1}| + |y_{i,t} - y_{i,t-1}| \leq 1, \quad \forall i \in \{1, \dots, N\}, \forall t \geq 1.$$

This constraint enforces the kinematic limits of the robots, ensuring that position changes between consecutive time steps are physically realizable. The Manhattan distance bound of 1 permits movement to any of the four adjacent cells or staying in place, but prohibits diagonal moves and larger jumps.

3. Energy budget:

$$D_{i,t} = \sum_{\tau=1}^t (|x_{i,\tau} - x_{i,\tau-1}| + |y_{i,\tau} - y_{i,\tau-1}|) \leq \ell, \quad \forall i \in \{1, \dots, N\}, \forall t \geq 0.$$

This constraint ensures that the total distance traveled by each robot at any time does not

exceed its battery capacity ℓ . Since all robots are identical, they share the same energy limit. This forces the planner to make strategic decisions about which regions to explore and which robots should undertake longer journeys versus remaining in supporting roles. Once a robot reaches its energy limit, it must remain stationary for the remainder of the mission.

4. Obstacle avoidance (with online discovery):

$$p_{i,t} \notin O_t, \quad \forall i \in \{1, \dots, N\}, \forall t \geq 0,$$

where $O_t \subseteq O_{t+1}$ for all $t \geq 0$.

These constraints prevent robots from planning to occupy cells containing known obstacles. Since obstacles are discovered gradually ($O_t \subseteq O_{t+1}$), the constraint set grows over time as exploration proceeds. This formulation naturally handles the online nature of the problem: robots must avoid all obstacles known at time t when planning their position at time t , but cannot anticipate obstacles that will only be discovered later. When a robot encounters a previously unknown obstacle during execution (triggering the obstacle encounter penalty P_{obstacle}), it uses reactive obstacle avoidance to navigate around the obstruction.

5. Collision avoidance:

$$p_{i,t} \neq p_{j,t}, \quad \forall i, j \in \{1, \dots, N\} \text{ with } i \neq j, \forall t \geq 0.$$

This constraint ensures that no two robots occupy the same cell simultaneously, preventing inter-robot collisions. Combined with the motion constraints, this also implicitly prevents certain types of edge collisions where robots would swap positions in a single time step.

6. Coverage update:

$$C_{x,y} \geq \mathbb{I}[p_{i,t} = (x, y)], \quad \forall x \in \{1, \dots, n\}, \forall y \in \{1, \dots, m\}, \forall i \in \{1, \dots, N\}, \forall t \geq 0.$$

These constraints link the coverage variables to the robot positions, ensuring that the coverage term in the objective function accurately reflects the set of visited cells. Since $C_{x,y}$ appears with a positive coefficient in the objective and we are maximizing, the optimizer will set $C_{x,y} = 1$ whenever any robot visits cell (x, y) at any time.

Decision Variables

The optimization determines the complete trajectory of each robot from the initial time until either all energy is exhausted or exploration is complete:

$$\{ p_{i,t} = (x_{i,t}, y_{i,t}) \mid i \in \{1, \dots, N\}, t \in \{0, 1, 2, \dots, T_{\max}\} \},$$

where $T_{\max} = \max\{\ell, nm\}$ is an upper bound on mission duration.

Given the initial positions $\{p_{i,0}\}_{i=1}^N$ and the current obstacle knowledge O_0 , the optimization must select positions for all robots at all future time steps that satisfy all constraints while maximizing the objective function. This results in a total of $2NT_{\max}$ integer decision variables (two coordinates per robot per time step), plus nm binary coverage variables, plus $NT_{\max}$ binary disconnection indicators $\delta_{i,t}$, plus $NT_{\max}$ binary obstacle encounter indicators $\eta_{i,t}$, making this a challenging mixed-integer nonlinear optimization problem due to the distance computations in the connectivity and penalty terms.

The planner generates the full path for each robot from start to finish in one optimization solve. This approach differs from receding horizon methods and provides globally optimal paths subject to the known information at planning time. However, since obstacles are discovered online, the actual executed paths may deviate from the plan when reactive obstacle avoidance is triggered. In practice, replanning may be performed periodically as new obstacle information becomes available, but each planning cycle produces a complete path rather than just the next few steps.

2.1 Simulated Annealing Implementation

The implemented **Simulated Annealing (SA)** algorithm optimizes the movement sequences of six robots over 125 time steps on a 20×20 grid. The algorithm begins from a feasible random path and iteratively explores new neighboring solutions by modifying up to 15 movements per robot. Each candidate path is validated against energy, collision, and obstacle constraints.

Solutions are evaluated using the following cost function:

$$\text{Cost} = \frac{\alpha}{\text{Coverage}} + \frac{\beta}{\text{Connectivity}} + \gamma \cdot P_{\text{disconnect}}$$

where $\alpha = 3000$, $\beta = 500$, and $\gamma = 4.0$ weight coverage, connectivity, and disconnection penalties respectively. The algorithm always accepts better solutions, while worse solutions are accepted with a probability based on the current temperature T . The temperature cools geometrically from 10.0 to 0.1 across 5000 iterations, gradually reducing exploration.

2.1.1 Pseudocode

Algorithm 1 Simulated Annealing for Multi-Robot Path Planning

```

1: Initialize current solution  $S_{current}$ 
2: Evaluate cost  $C_{current} = f(S_{current})$ 
3: Set best solution  $S_{best} \leftarrow S_{current}$  and  $C_{best} \leftarrow C_{current}$ 
4: Initialize temperature  $T \leftarrow T_{init}$ 
5: Set iteration counter  $iter \leftarrow 0$ 
6: while  $T > T_{min}$  and  $iter < N_{max}$  do
7:   Generate neighbor solution  $S_{new}$  from  $S_{current}$ 
8:   if  $S_{new}$  is feasible then
9:     Compute cost  $C_{new} = f(S_{new})$ 
10:    if  $C_{new} < C_{current}$  then
11:      Accept  $S_{new}$  as  $S_{current}$ 
12:    else
13:      Compute acceptance probability  $p \leftarrow \exp\left(-\frac{C_{new} - C_{current}}{T}\right)$ 
14:      Generate random number  $r \in [0, 1]$ 
15:      if  $r < p$  then
16:        Accept  $S_{new}$  as  $S_{current}$ 
17:      else
18:        Reject  $S_{new}$ 
19:      end if
20:    end if
21:    if  $C_{new} < C_{best}$  then
22:      Update global best:  $S_{best} \leftarrow S_{new}$ ,  $C_{best} \leftarrow C_{new}$ 
23:    end if
24:  end if
25:  Update temperature:  $T \leftarrow \alpha \times T$  ▷ Geometric cooling
26:  Increment  $iter$ 
27: end while
28: return  $S_{best}$ 

```

2.1.2 Flowchart




Figure 2.2: Simulated Annealing algorithm flowchart for multi-robot path planning.

2.1.3 Problem-Specific Considerations

- **Neighborhood Operator:** Implemented in `generate_neighbor()`, which randomly alters up to 15 movements per robot and ensures feasibility before returning the modified path.
- **Feasibility Check:** Conducted by `is_feasible(path_array)`, verifying that paths respect motion constraints, avoid collisions, remain within map bounds, and comply with energy limits.
- **Cooling Schedule:** The temperature T is reduced geometrically using $T_{new} = T_{current} \times 0.995$, gradually limiting the acceptance of worse solutions as the algorithm progresses.

2.1.4 GenerateNeighbor Subroutine

Algorithm 2 Generate Feasible Neighbor Solution

Require: Current movements $M_{current}$, maximum retries R_{max}

Ensure: Feasible neighbor movements $M_{neighbor}$

```

1: for  $attempt \leftarrow 1$  to  $R_{max}$  do
2:    $M_{new} \leftarrow \text{Copy}(M_{current})$  ▷ Deep copy
3:   for each robot  $r$  do
4:     for  $i \leftarrow 1$  to 15 do ▷ Modify 15 random movements
5:        $idx \leftarrow \text{RandomInt}(0, K - 1)$  ▷ Random step index
6:        $M_{new}[r][idx] \leftarrow \text{RandomInt}(0, 4)$  ▷ Random direction: 0=stay, 1-4=move
7:     end for
8:   end for
9:    $P_{new} \leftarrow \text{MovementsToPositions}(M_{new})$ 
10:  if  $\text{IsFeasible}(P_{new})$  then ▷ Check for collisions
11:    return  $M_{new}$  ▷ Valid neighbor found
12:  end if
13: end for
14: Warning: No feasible neighbor found
15: return  $M_{current}$  ▷ Return original if all attempts fail

```

2.1.5 Key Algorithm Components

Cost Function

The cost function is designed as a **minimization problem**:

$$\text{Cost}(P) = -\alpha \cdot \text{Coverage}(P) + \beta \cdot \text{PathLength}(P) \quad (2.1)$$

where:

- $\text{Coverage}(P)$: Percentage of area covered by robots
- $\text{PathLength}(P)$: Total path length traveled
- α : Weight for coverage (negative to maximize coverage)
- β : Weight for path smoothness

Lower cost is better - the algorithm minimizes this cost function.

Acceptance Probability

The Metropolis criterion for accepting worse solutions:

$$P(\text{accept}) = \begin{cases} 1 & \text{if } C_{\text{new}} < C_{\text{current}} \\ \exp\left(-\frac{C_{\text{new}} - C_{\text{current}}}{T}\right) & \text{otherwise} \end{cases} \quad (2.2)$$

As temperature T decreases:

- Early iterations: High T allows exploration (accepts many worse solutions)
- Late iterations: Low T focuses on exploitation (rarely accepts worse solutions)

Cooling Schedule

Geometric cooling is used:

$$T_{k+1} = \alpha \cdot T_k \quad (2.3)$$

where α is the cooling rate (typically 0.995 in your implementation).

Neighbor Generation Strategy

For each robot, 15 random movement steps are modified:

- Movement values: 0 (stay), 1 (up), 2 (down), 3 (left), 4 (right)
- Feasibility check: Ensures no robot collisions
- Maximum 100 retries to find valid neighbor

Termination Criteria

The algorithm stops when:

1. Temperature drops below $T_{\min}$, **OR**
2. Maximum iterations $N_{\max}$ is reached

2.2 Genetic Algorithm Implementation

The implemented **Genetic Algorithm (GA)** algorithm optimizes the movement sequences of six robots over 125 time steps on a 20×20 grid. The algorithm begins from a feasible random population, and then does several generational updates by mutation and crossover methods searching the solution landscape until it reaches the optimal solution. Each candidate path is validated against energy, collision, and obstacle constraints.

Solutions are evaluated using the following cost function:

$$\text{Cost} = \frac{\alpha}{\text{Coverage}} + \frac{\beta}{\text{Connectivity}} + \gamma \cdot P_{\text{disconnect}}$$

where $\alpha = 3000$, $\beta = 500$, and $\gamma = 4.0$ weight coverage, connectivity, and disconnection penalties respectively.

The algorithm chooses a certain percentage of chromosomes (solutions) per population as elites, where they are propagated to the next iteration, since they have the best fitness value (lowest cost). It selects another percentage of individuals to mutate, and those are the worst fitness individuals. Then, parents are selected from the remaining solutions to perform crossover.

2.2.1 Pseudocode

Algorithm 3 Genetic Algorithm for Multi-Robot Path Planning

```

1: Input: Initial solution, Robot positions, Population size  $m$ , Generation size  $i_{max}$ 
2: Input: Elite rate  $p_e$ , Mutation rate  $p_m$ , Objective function  $f(\mathbf{x})$ 
3: Initialize: Generate initial population from initial solution
4: Initialize: Best solution  $\mathbf{x}^*$ , Best cost  $f^* = \infty$ 
5: while  $i \leq i_{max}$  do
6:   // Evaluate Fitness
7:   for each individual in population do
8:     Compute fitness:  $fitness = f(\mathbf{x})$ 
9:     Store fitness in  $fitness\_list$ 
10:  end for
11:  // Update Best Solution
12:   $current\_cost = \min(fitness\_list)$ 
13:   $current\_solution = population[argmin(fitness\_list)]$ 
14:  if  $current\_cost < f^*$  then
15:     $f^* = current\_cost$ 
16:     $\mathbf{x}^* = current\_solution$ 
17:  end if
18:  // Generate New Population
19:   $new\_population = []$ 
20:  // Elitism: Select top  $p_e \times m$  survivors
21:  Sort population by fitness
22:   $elites = \text{top } p_e \times m \text{ individuals}$ 
23:  Add  $elites$  to  $new\_population$ 
24:  // Mutation: Mutate worst  $p_m \times m$  individuals
25:   $worst\_individuals = \text{bottom } p_m \times m \text{ individuals}$ 
26:  for each individual in  $worst\_individuals$  do
27:    Apply swap mutation per robot path
28:    Ensure mutated individual is feasible
29:    Add mutated individual to  $new\_population$ 
30:  end for
31:  // Crossover: Generate offspring for remaining slots
32:   $num\_crossover = m - |new\_population|$ 
33:   $offspring\_count = 0$ 
34:  while  $offspring\_count < num\_crossover$  do
35:    Select 2 parents using Stochastic Universal Sampling (SUS)
36:    Apply one-point crossover per robot path
37:    Generate  $offspring_1, offspring_2$ 
38:    Filter feasible offspring
39:    Sort feasible offspring by fitness
40:    if feasible offspring exist then
41:      Add fittest offspring to  $new\_population$ 
42:       $offspring\_count = offspring\_count + 1$ 
43:    end if
44:  end while

```

2.2.2 Problem-Specific Considerations

- **Crossover Operator:** Implemented using one-point crossover per robot path in `one_point_crossover_for_robots_paths()`, which selects a random crossover point for each robot's path independently and exchanges segments between two parent solutions to create offspring.
- **Mutation Operator:** Implemented using swap mutation per robot path in `swap_mutation_per_robot_path()`, which randomly swaps two positions within each robot's path independently to introduce variation.
- **Parent Selection:** Implemented using Stochastic Universal Sampling (SUS) in `select_parents_by_sus()`, which uses evenly-spaced pointers to select parents proportionally to their fitness, providing better diversity than roulette wheel selection.
- **Feasibility Check:** Conducted by `is_feasible(path_array)` throughout the genetic operations, verifying that generated offspring and mutated individuals respect motion constraints, avoid collisions, remain within map bounds, and comply with energy limits. Infeasible solutions are rejected and regenerated.
- **Elite Preservation:** The top $p_e \times m$ individuals (fittest solutions) are directly carried over to the next generation to preserve high-quality solutions.

2.3 Ant Colony Optimization Implementation

The **Ant Colony Optimization (ACO)** module constructs paths for all robots by sampling movements for a population of ants while maintaining a global pheromone map. Each ant independently samples a path using either a pheromone-only strategy (SACO) or the Ant System variant (AS) that mixes pheromone strength with a heuristic map. After every iteration, pheromones evaporate globally and are reinforced along the best path of the iteration; local pheromone updates occur either per move (AS) or once per constructed path (SACO). Heuristic desirability is pre-computed so that unexplored cells receive higher weight than free space and obstacles.

2.3.1 Pseudocode

Algorithm 4 Ant Colony Optimization for Multi-Robot Path Planning

```

1: Input: Number of ants  $m$ , iterations  $N_{max}$ ,  $\alpha_\tau$  (pheromone weight),  $\beta_\eta$  (heuristic weight),
   evaporation rate  $\rho$ , strategy  $\in \{\text{SACO}, \text{AS}\}$ 
2: Initialize pheromone map  $\tau$  to 1.0 for all cells; build heuristic map  $\eta$  (unexplored > free >
   obstacle)
3: Initialize:  $best\_cost \leftarrow \infty$ ,  $best\_path \leftarrow \emptyset$ 
4: for  $iter = 1$  to  $N_{max}$  do
5:   Apply evaporation:  $\tau \leftarrow (1 - \rho) \cdot \tau$ 
6:   for ant  $k = 1$  to  $m$  do
7:     for each robot  $r$  do
8:        $p \leftarrow$  initial position of robot  $r$ 
9:       for  $t = 1$  to  $K$  steps do
10:         $F \leftarrow$  feasible moves from  $p$  (stay or 4-connected within bounds)
11:        Compute weight  $w = \tau^{\alpha_\tau} \cdot \eta^{\beta_\eta}$  for each move (omit  $\eta$  when SACO)
12:        Sample move via roulette-wheel on  $w$ 
13:        Apply move, update  $p$ ; if AS, apply local decay on  $\tau$  at  $p$ 
14:      end for
15:    end for
16:    Build full path  $P_k$ , compute cost  $C_k$  with coverage/connectivity penalties
17:    Local deposit along  $P_k$  (pheromone-only or heuristic-weighted)
18:  end for
19:  Select iteration best  $(P^*, C^*)$ ; reinforce  $\tau$  along  $P^*$  (AS) while evaporating
20:  if  $C^* < best\_cost$  then
21:     $best\_cost \leftarrow C^*$ ,  $best\_path \leftarrow P^*$ 
22:  end if
23:  Record metrics / visualization
24: end for
25: Return  $best\_path, best\_cost$ 

```

2.3.2 Implementation Details

- **Move sampling:** `AntColonyStrategy.select_move()` computes weights from τ^{α_τ} (SACO) or $\tau^{\alpha_\tau} \eta^{\beta_\eta}$ (AS). Moves are drawn with roulette-wheel sampling; if all weights vanish, a random feasible move is chosen.
- **Heuristic map:** `_build_heuristic_map()` assigns 2.0 to unexplored, 1.2 to free space, and 10^{-6} to obstacles to discourage infeasible cells.
- **Evaporation and updates:** `before_iteration()` applies uniform evaporation. `ASStrategy.local_decay()`

applies per-move decay toward the initial pheromone. Both strategies call `update_local_pheromone` after constructing a full path, depositing $(1/(1+\text{cost}))$ on visited cells (scaled by heuristic for AS).

- **Global reinforcement:** `ASStrategy.after_iteration()` reinforces the iteration-best path with factor $(1 + \beta_\eta)/(1 + \text{cost})$ while re-applying evaporation; SACO relies solely on local deposits plus the global evaporation step.
- **Feasibility:** `_get_feasible_moves()` restricts motion to in-bounds 4-connected moves; collision/energy/obstacle feasibility is enforced in the shared cost function.
- **Interfaces:** Inherits visualization/cost utilities from `BaseOptimizer`; maintains ant moves as an $(m \times R \times K)$ array and converts to positions with `movements_to_positions()` for scoring.

2.4 Artificial Bee Colony Implementation

The **Artificial Bee Colony (ABC)** module optimizes the same movement-matrix representation using employed, onlooker, and scout phases. Each food source is a full set of robot trajectories (robots $\times$ horizon); feasibility is enforced via the shared constraints (bounds, 4-connected motion, energy budget, obstacle avoidance, no same-cell collisions). Fitness is $1/(\text{cost} + \epsilon)$ where cost uses the common objective $(\alpha/\text{coverage} + \beta/\text{connectivity} + \gamma \cdot P_{\text{disconnect}})$. Infeasible neighbors are rejected so the colony always evolves within the feasible region.

2.4.1 Algorithm Flow

- **Initialization:** Build a feasible initial colony (default 24 sources) using the random feasible path generator; track a trial counter per source and the global best.
- **Employed phase:** One employed bee per source. Pick a random peer and perturb a short window of one robot's path by mixing peer moves with random moves; occasionally swap two moves in another robot path for diversity. Greedy acceptance on improved cost; otherwise increment that source's trial counter.
- **Onlooker phase:** Allocate onlookers proportionally to fitness (roulette). Each onlooker picks a source, applies the same neighbor operator with a random peer, and greedily accepts if better; otherwise increments that source's trial counter.
- **Scout phase:** Any source whose trial counter exceeds a limit (default 12) is abandoned and replaced with a fresh feasible random path; its trial counter resets to zero. This injects exploration when a source stagnates.

- **Termination:** Run for a fixed number of cycles (default 200) or until a user-defined convergence/stall condition; report the best source and cost.

2.4.2 Key Operators and Parameters

- **Neighbor perturbation:** Window length 10 steps by default; blend peer moves with random moves (60% peer bias inside the window) and allow a 30% chance of an intra-path swap on another robot to adjust connectivity. Up to 5 feasible-attempt retries; fall back to the original source if none are feasible.
- **Selection and replacement:** Greedy acceptance in employed/onlooker phases; scouts reinitialize stagnant sources after hitting the trial limit.
- **Hyperparameters:** Colony size 24, cycles 200, limit 12, onlooker ratio 0.5 (half bees act as onlookers per cycle), neighbor window 10, neighbor attempts 5. Case studies inherit these defaults unless overridden in `case_study_config.py`.
- **Integration:** Uses the shared cost function and feasibility checks; visualization hooks mirror GA/SA for live or replay modes.

2.4.3 Pseudocode

Algorithm 5 Artificial Bee Colony for Multi-Robot Path Planning

```

1: Input: Colony size  $m$ , cycles  $N_{max}$ , trial limit  $L$ , onlooker ratio  $\rho$ , neighbor window  $w$ ,
   attempts  $a$ 
2: Initialize  $m$  feasible sources (movement matrices) and trial counters  $trials_i \leftarrow 0$ 
3: Evaluate all sources, set  $best\_cost \leftarrow \infty$ 
4: for  $cycle = 1$  to  $N_{max}$  do
5:   // Employed phase
6:   for each source  $i$  do
7:     pick peer  $k \neq i$  uniformly
8:     neighbor  $\leftarrow$  perturb_window( $source_i$ ,  $source_k$ ,  $w$ ,  $a$ )
9:     if feasible(neighbor) and cost(neighbor) < cost( $source_i$ ) then
10:       $source_i \leftarrow$  neighbor;  $trials_i \leftarrow 0$ 
11:     else
12:       $trials_i++$ 
13:     end if
14:   end for
15:   // Onlooker phase
16:   compute fitness  $fit_i = 1/(cost_i + \varepsilon)$ , probabilities  $p_i = fit_i / \sum fit$ 
17:   for  $j = 1$  to  $\lfloor \rho m \rfloor$  do
18:     sample source index  $i \sim p$ , peer  $k \neq i$ 
19:     neighbor  $\leftarrow$  perturb_window( $source_i$ ,  $source_k$ ,  $w$ ,  $a$ )
20:     if feasible(neighbor) and cost(neighbor) < cost( $source_i$ ) then
21:       $source_i \leftarrow$  neighbor;  $trials_i \leftarrow 0$ 
22:     else
23:       $trials_i++$ 
24:     end if
25:   end for
26:   // Scout phase
27:   for each source  $i$  do
28:     if  $trials_i \geq L$  then
29:       $source_i \leftarrow$  new_feasible_random();  $trials_i \leftarrow 0$ 
30:     end if
31:   end for
32:   update  $best\_cost$ ,  $best\_source$  if any source improved the global best
33:   record metrics / visualization
34: end for
35: Return best source and cost

```

Chapter 3

Results

3.1 Simulated Annealing Results

3.1.1 Parameters Tuning

The following parameters were tuned for different case studies to analyze algorithm performance under various conditions.

Table 3.1: Performance Metrics: Area Coverage Results Summary

Case	Robots	Steps	α	β	Temp	Iterations	Coverage (%)
C1	6	125	3000	500	10.0	910	69.2
C2	6	125	400	50	10.0	910	63.0
C3	6	125	1000	200	10.0	910	54.2
C4	6	125	3000	500	100.0	1370	55.0
C5	6	125	3000	500	500.0	1700	55.2
C6	6	125	400	50	100.0	1370	61.8

Discussion of Results: The experimental cases (C1–C6) demonstrate how varying the cost function weights (α , β) and temperature parameters influences the optimization outcomes.

- **Case C1:** This baseline configuration ($\alpha = 3000$, $\beta = 500$, Temp = 10.0) achieved the highest coverage (69.2%). The high α weight strongly encourages spatial dispersion, ensuring that robots spread efficiently to unexplored regions, while the moderate β maintains sufficient communication links to prevent disconnection. The moderate temperature allows early-stage exploration followed by stable convergence, making this configuration the most balanced overall.

- **Case C2:** With reduced weights ($\alpha = 400$, $\beta = 50$), the optimization placed less emphasis on both coverage and connectivity. As a result, the robots explored less aggressively, lowering total coverage (63.0%). The algorithm became less sensitive to cost differences, leading to fewer beneficial path updates.
- **Case C3:** This setup ($\alpha = 1000$, $\beta = 200$) was designed to test moderate weighting. However, the reduced α compared to C1 decreased the motivation for robots to cover new ground. The algorithm converged faster but prematurely, limiting exploration diversity and resulting in a further drop in coverage (54.2%). This shows that smaller α values increase convergence speed but sacrifice global search performance.
- **Case C4:** By increasing the initial temperature to 100.0 (while keeping $\alpha = 3000$, $\beta = 500$), the algorithm allowed more random acceptance of worse solutions early on. Although this expanded exploration, it also delayed convergence and caused some oscillation in solution quality. The final coverage (55.0%) dropped because the system spent too long exploring suboptimal regions before stabilizing. This demonstrates that excessively high initial temperatures can hinder focused improvement.
- **Case C5:** Here, the initial temperature was increased even further to 500.0. The system's acceptance probability for worse solutions became extremely high in early iterations, making the search nearly random. Despite having the same weights as C1, the algorithm required more iterations (1700) to cool down and stabilize, yet only reached 55.2% coverage. This confirms that while high temperature promotes exploration, it can significantly degrade convergence efficiency if not controlled by an appropriate cooling rate.
- **Case C6:** This configuration combined lower weights ($\alpha = 400$, $\beta = 50$) with a high temperature (100.0). The weak weighting reduced the importance of both coverage and connectivity, while the high temperature further amplified random exploration. Interestingly, the result (61.8%) was slightly better than C3 and C4 because the higher temperature compensated partially for low weights by encouraging movement diversity, but the algorithm still lacked the focused improvement of C1.

Overall, these cases illustrate that moderate temperature and balanced cost weights yield the best results. When α and β are too low, the optimization loses direction, while excessively high temperature values disrupt convergence. The best trade-off between exploration and stability is achieved in C1, where coverage and connectivity are well-balanced and the system converges efficiently.

3.1.2 Cooperative Robot Number Configurations

Table 3.2: Performance Metrics for Different Robot Counts

Case	Robots	Steps	α	β	Temp	Iterations	Coverage (%)
C7	4	130	3000	500	10.0	910	52.0
C8	8	90	3000	500	10.0	910	66.5

Discussion of Results: The results from cases C7 and C8 highlight the influence of robot team size on overall area coverage and coordination efficiency under identical optimization conditions.

- **Case C7 (4 Robots):** With only four robots exploring the 20×20 grid, the available sensing and movement capacity is significantly reduced, resulting in a coverage of 52.0%. Although each robot has more space to explore, the reduced team size limits simultaneous area exploration and increases travel distance per robot. Moreover, maintaining connectivity between fewer agents restricts how far each one can move independently, further constraining exploration. This case demonstrates that a small team suffers from both limited parallelism and higher coordination burden per agent.
- **Case C8 (8 Robots):** Increasing the team to eight robots improves spatial distribution and parallel exploration, leading to higher coverage (66.5%). The additional robots reduce overlap in exploration paths and allow better partitioning of the map. Connectivity is also easier to maintain since robots can stay within range of at least one neighbor while still dispersing effectively. Compared to C7, the system shows higher cooperative efficiency and better coverage saturation, proving that a moderate increase in robot count enhances global performance without overwhelming the optimization with excessive interactions.

Overall, these cases confirm that expanding the robot team improves total coverage up to a point. However, as the number of robots grows, maintaining communication and avoiding path overlap become increasingly important to preserve efficiency. The results indicate that around 8–10 robots provide a good balance between cooperative reach and manageable coordination in the 20×20 environment.

3.1.3 Visualization

The visualization component provides real-time feedback on how the Simulated Annealing (SA) algorithm explores the solution space and how the multi-robot paths evolve over time.

Path and Coverage Map Visualization

This visualization shows the environment map, the planned paths of all robots, and the area they have collectively covered. It helps to visually confirm that constraints (like obstacle avoidance and no-collision) are met and provides an intuitive measure of the solution's objective (maximizing coverage).




Figure 3.1: Graphical visualization of the optimization process, showing robot paths and the area covered in the environment.

Terminal and Metric Visualization

In addition to the graphical output, the terminal provides essential diagnostic information. This output is critical for tracking the results of the SA algorithm.




Figure 3.2: Terminal visualization and coverage results, displaying key performance metrics.

3.1.4 Conclusions

This study implemented a **Simulated Annealing** approach for multi-robot area coverage optimization in a 20×20 search-and-rescue grid.

Key Observations

- **Best Performance:** Case C1 achieved the highest coverage (69.2%) using balanced weights and moderate temperature.
- **Temperature Sensitivity:** Excessively high initial temperatures (C4, C5) reduced convergence reliability.
- **Weight Impact:** Higher α and β values strengthened area coverage but slightly increased runtime.
- **Robot Cooperation:** Increasing the number of robots improved total coverage, confirming cooperative exploration benefits.

Parameter Sensitivity

- Coverage performance is highly sensitive to the relative weighting between α and β .
- Moderate initial temperature (10.0–20.0) provides the best convergence–exploration balance.
- Increasing robot count consistently improves coverage but introduces additional coordination challenges.

Overall, the study shows that careful tuning of these hyperparameters is essential to maximize multi-robot exploration efficiency using simulated annealing.

3.2 Genetic Algorithm Results

3.2.1 Parameters Tuning

The following parameters were tuned for different case studies to analyze algorithm performance under various conditions.

Table 3.3: Performance Metrics: Area Coverage Results Summary for mutation rate = 30% and elite rate 10% while varying Population and Generation Size

Case	Robots	Steps	α	β	Population	Generation	Coverage (%)
C1	6	125	3000	500	20	200	76.2
C2	6	125	3000	500	10	200	65.5
C3	6	125	3000	500	30	200	71.0
C4	6	125	3000	500	20	100	68.5
C5	6	125	3000	500	20	300	70.8
C6	6	125	3000	500	20	400	73.2

Table 3.4: Performance Metrics: Area Coverage Results Summary for mutation rate = 30% and elite rate 10% while varying α and β

Case	Robots	Steps	α	β	Population	Generation	Coverage (%)
C7	6	125	100	200	20	200	51.0
C8	6	125	9000	7000	20	200	59.5
C9	6	125	10,000	670	20	200	69.8

Table 3.5: Performance Metrics: Area Coverage Results Summary for Population size = 20 and Generation Size 200.

Case	Robots	Steps	α	β	Mutation Rate	Elitism Rate	Coverage (%)
C10	6	125	3000	500	0.1	0.1	74.8
C11	6	125	3000	500	0.1	0.4	68.0
C12	6	125	3000	500	0.4	0.1	65.2

Discussion of Results: The experimental cases demonstrate how varying population size, generation count, cost function weights (α , β), mutation rate, and elitism rate influence the genetic algorithm’s optimization performance.

- **Cases C1–C6 (Population and Generation Size Variation):** Case C1 achieved the highest coverage (76.2%) with a population size of 20 and 200 generations, representing the optimal balance between genetic diversity and computational efficiency. Reducing the population to 10 (C2) significantly decreased coverage to 65.5%, as the limited genetic diversity restricted the algorithm’s ability to explore the solution space effectively and escape local optima. Increasing the population to 30 (C3) yielded 71.0% coverage, slightly worse than C1, suggesting diminishing returns from larger populations due to increased computational overhead per generation without proportional improvement in solution quality. Varying generation count showed that 100 generations (C4, 68.5%) was insufficient for convergence, while 200 generations (C1, 76.2%) provided adequate time for crossover and selection to refine solutions. Further increases to 300 (C5, 70.8%) and 400 (C6, 73.2%) generations showed modest improvements, indicating near-convergence by generation 200.
- **Cases C7–C9 (Cost Function Weight Variation):** These cases examined the impact of α (coverage weight) and β (connectivity weight) on optimization outcomes. Case C7 with low weights ($\alpha = 100$, $\beta = 200$) achieved only 51.0% coverage, as the weak cost signal failed to adequately penalize poor coverage and guide the genetic operators toward better solutions. Case C8 with very high weights ($\alpha = 9000$, $\beta = 7000$) reached 59.5% coverage, performing worse than expected. The excessively large weights created a steep fitness landscape that made small improvements difficult to detect, reducing the effectiveness of parent selection (SUS) in distinguishing between solutions and potentially causing premature convergence. Case C9 with $\alpha = 10000$ and $\beta = 670$ achieved 69.8% coverage, demonstrating that heavily prioritizing coverage (α) over connectivity (β) improved exploration. However, it still underperformed compared to C1’s balanced weighting ($\alpha = 3000$, $\beta = 500$), suggesting that some connectivity maintenance is necessary for coordinated multi-robot coverage.
- **Cases C10–C12 (Mutation and Elitism Rate Variation):** Case C10 with 10% muta-

tion rate and 10% elitism rate achieved 74.8% coverage, nearly matching the best result from C1. The low mutation rate (10%) applied swap mutations sparingly to the worst individuals, preserving good genetic material while introducing controlled variation. The low elitism rate (10%) ensured only the top 2 individuals (from population of 20) survived unchanged, balancing preservation of elite solutions with allowing new genetic combinations through crossover. Increasing elitism to 40% (C11, 68.0%) significantly reduced coverage, as preserving 8 elite individuals per generation limited the diversity of new offspring generated through crossover. This demonstrates that excessive elitism can cause premature convergence by restricting the gene pool. Conversely, increasing mutation rate to 40% (C12, 65.2%) also decreased performance, as aggressively mutating 8 of the worst individuals per generation introduced too much randomness, disrupting the accumulation of beneficial genetic patterns established through crossover and selection.

Overall, these results demonstrate that the genetic algorithm performs optimally with moderate hyperparameters: a population size of 20 provides sufficient genetic diversity without excessive computational cost, 200 generations allow adequate convergence time, balanced cost weights ($\alpha = 3000$, $\beta = 500$) create appropriate fitness gradients for selection, and low mutation (10%) and elitism (10%) rates balance exploration with exploitation. The combination of one-point crossover per robot path, SUS parent selection, and swap mutation per robot path effectively explores the multi-robot path planning solution space when properly tuned.

3.2.2 Cooperative Robot Number Configurations

Table 3.6: Performance Metrics for Different Robot Counts

Case	Robots	Steps	α	β	Population Size	Generation Size	Coverage (%)
C13	4	130	3000	500	20	200	55.2
C14	8	90	3000	500	20	200	68.0

Discussion of Results: The results from cases C13 and C14 highlight the influence of robot team size on overall area coverage and coordination efficiency under identical optimization conditions.

- **Case C13 (4 Robots):** With only four robots exploring the 20×20 grid over 130 steps, the coverage achieved 55.2%. Despite having more steps per robot (130 vs. 90 in C14), the limited team size fundamentally constrains parallel exploration capability. The genetic algorithm's crossover operator (one-point crossover per robot path) has fewer robot paths to recombine, reducing the diversity of spatial coverage patterns that can emerge through genetic recombination. Additionally, the smaller population of feasible solutions in the initial generation limits the genetic material available for SUS parent selec-

tion, potentially causing the algorithm to converge to suboptimal coverage distributions. The connectivity constraint ($\beta = 500$) further restricts exploration, as maintaining communication links between only four robots requires them to remain relatively clustered, limiting their ability to disperse across the entire map. This case demonstrates that a small robot team reduces both the solution space diversity for genetic operations and the physical capacity for simultaneous area exploration.

- **Case C14 (8 Robots):** Increasing the team to eight robots improves coverage to 68.0%, despite each robot taking fewer steps (90 vs. 130 in C13). The larger team provides the genetic algorithm with more dimensions to optimize (8 robot paths vs. 4), allowing crossover and mutation operators to explore a richer solution space. With eight robots, the one-point crossover per robot path can generate more diverse offspring by mixing different combinations of robot trajectories, enabling better spatial distribution patterns to emerge through selection. The SUS parent selection mechanism benefits from having solutions with more varied coverage patterns, as the fitness landscape becomes more nuanced when evaluating how eight robots collectively cover the map. The connectivity constraint is easier to satisfy with more robots, as each robot has more potential neighbors within communication range, allowing greater individual exploration freedom while maintaining the connected network. Compared to C13, the system demonstrates that the genetic algorithm's performance improves with larger teams because the increased dimensionality of the solution space provides more opportunities for beneficial genetic recombinations and mutations to discover effective multi-robot coverage strategies.

Overall, these cases confirm that expanding the robot team from 4 to 8 significantly improves coverage (55.2% to 68.0%) by enriching the solution space available for genetic operations. The genetic algorithm's crossover, mutation, and selection operators become more effective with larger teams because they can exploit more complex path combinations and spatial distributions. However, as the number of robots grows further, the solution space dimensionality increases exponentially, potentially requiring larger population sizes or more generations to adequately explore. The results suggest that 6–8 robots provide a good balance for the genetic algorithm with population size 20 and 200 generations, as they offer sufficient solution complexity for effective genetic search while remaining computationally manageable.

3.2.3 Testing Crossover, Mutation and Parent Selection Methods

Table 3.7: Here we use all the parameters same as C1 but vary the Mutation, Crossover and Parent Selection Methods.

Case	mutation_method	crossover_method	parent_selection_method	Coverage
C15	"swap"	"one_point"	"roulette"	47.0%
C16	"swap_per_robot_path"	"one_point"	"sus"	62.2%
C17	"swap"	"one_point_per_robots_paths"	"roulette"	55.5%
C18	"swap_per_robot_path"	"one_point_per_robots_paths"	"roulette"	62.5%

To specify those methods, while running the optimization, you pass in the following variables:

The configuration used for C1:

```
best_movements, best_cost = ga.run(
    initial_movements, # related to algorithm
    mutation_method="swap_per_robot_path",
    parent_selection_method="sus",
    crossover_method="one_point_per_robots_paths",
    robot_positions=ROBOTS_POSITIONS # related to algorithm
)
```

Table 3.8: Genetic Operator Methods Summary

Method	Description
<i>Mutation Methods</i>	
swap	Swaps the entire robots paths with one another.
swap_per_robot_path	Independently swaps two positions within each robot's path, preserving robot boundaries. Maintains path structure and significantly increases feasibility of mutated solutions.
<i>Crossover Methods</i>	
one_point	Applies crossover across entire robot paths, rather than within each movement step for each robots paths.
one_point_per_robot	Applies independent one-point crossover to each robot's path with separate crossover points. Preserves robot path structure while enabling effective genetic recombination.
<i>Parent Selection Methods</i>	
roulette	Fitness-proportional selection where each parent is selected independently with probability proportional to fitness. Higher variance can lead to premature convergence by over-selecting highly fit individuals.
sus	Stochastic Universal Sampling uses evenly-spaced pointers for uniform population sampling. Maintains better genetic diversity and prevents premature convergence to local optima.

Discussion of Results: These cases (C15–C17) compare different combinations of mutation, crossover, and parent selection methods to evaluate their impact on the genetic algorithm's performance, with all other parameters held constant at C1 values (population=20, generations=200, $\alpha = 3000$, $\beta = 500$).

- **Case C15 (swap + one_point + roulette):** This configuration achieved the lowest coverage (47.0%), using swap mutation on the entire solution array, standard one-point crossover across all robots simultaneously, and roulette wheel parent selection. The swap mutation operates on the flattened solution structure, randomly swapping movements between different robots' paths, which frequently violates the structure of individual robot trajectories and generates infeasible solutions that must be rejected. The one-point crossover method splits the entire solution at a single point, mixing the latter portions of all robot paths simultaneously. This destroys the spatial coherence of multi-

robot coordination patterns, as it doesn't respect the natural boundaries between individual robot paths. Roulette wheel selection, while providing fitness-proportional selection, has higher variance than SUS and can lead to premature convergence by repeatedly selecting the same high-fitness individuals, reducing genetic diversity. The combination of these methods creates a system where most genetic operations produce infeasible offspring, forcing the algorithm to repeatedly regenerate solutions and severely limiting effective exploration.

- **Case C16 (swap_per_robot_path + one_point + sus):** This configuration achieved 62.2% coverage, using swap mutation per robot path, standard one-point crossover, and SUS parent selection. The swap_per_robot_path mutation operates on each robot's path independently, swapping two positions within a single robot's movement sequence. This preserves the structure of other robots' paths and significantly increases the likelihood of generating feasible solutions. However, the standard one-point crossover still splits the entire solution array at a single point, mixing movements across robot boundaries. While less disruptive than in C15 (due to better mutation), it still occasionally breaks the coordination between robots. SUS parent selection provides more uniform sampling of the population compared to roulette, ensuring better genetic diversity and preventing premature convergence to local optima. The improved performance over C15 demonstrates that respecting robot path boundaries during mutation is critical, though the crossover method remains suboptimal.
- **Case C17 (swap + one_point_per_robots_paths + roulette):** This configuration achieved 55.5% coverage, using standard swap mutation, one-point crossover per robot path, and roulette wheel selection. The one_point_per_robots_paths crossover applies a separate crossover point to each robot's path independently, preserving the structure of individual robot trajectories while allowing beneficial path segments to be inherited from parents. This generates far more feasible offspring than standard one-point crossover. However, the standard swap mutation still operates on the flattened solution array, frequently swapping movements between different robots and creating infeasible solutions. Roulette wheel selection adds selection pressure variance, though its impact is less pronounced than the structural issues caused by swap mutation. The moderate performance improvement over C15 (47.0% to 55.5%) shows that structure-preserving crossover is beneficial, but the gains are limited by the disruptive mutation operator.
- **Case C18 (swap_per_robot_path + one_point_per_robots_paths + roulette):** This configuration achieved 62.5% coverage, using both structure-aware genetic operators (swap per robot path and one-point crossover per robot path) but with roulette wheel parent selection instead of SUS. The swap_per_robot_path mutation and one_point_per_robots_paths crossover both respect robot path boundaries, ensuring high feasibility rates for generated offspring and preserving beneficial multi-robot coordina-

tion patterns. However, roulette wheel selection introduces higher variance in parent selection compared to SUS. While fitness-proportional, roulette can select the same high-fitness individuals multiple times in a single generation, reducing genetic diversity and potentially causing premature convergence. The performance (62.5%) is significantly better than cases using structure-violating operators (C15: 47.0%, C17: 55.5%) and nearly matches C16 (62.2%), confirming that structure-aware mutation and crossover are the primary drivers of performance. However, C18 still underperforms C1 (76.2%) by 13.7 percentage points, demonstrating that parent selection method has measurable impact on final solution quality. The comparison between C18 (roulette) and C1 (SUS) isolates the effect of selection method, showing that SUS's more uniform sampling maintains better population diversity throughout evolution, enabling the genetic algorithm to escape local optima and continue improving for more generations.

- **Case C1 (optimal: swap_per_robot_path + one_point_per_robots_paths + sus):** For reference, C1 achieved 76.2% coverage using the optimal combination of all three operators. The swap_per_robot_path mutation maintains feasibility by operating within individual robot paths. The one_point_per_robots_paths crossover respects robot path boundaries, enabling effective recombination of successful trajectory patterns. SUS parent selection ensures consistent genetic diversity throughout evolution. This combination creates a synergistic effect where genetic operations consistently produce feasible, high-quality offspring that can be effectively evaluated and selected.

Overall, these results demonstrate that operator design must respect the problem structure for multi-robot path planning. Methods that treat the solution as a flat array (swap, one_point) frequently violate robot path boundaries, not explore enough of the solution space, generating infeasible solutions and wasting computational resources on rejected offspring. In contrast, structure-aware operators (swap_per_robot_path, one_point_per_robots_paths) preserve the independence of robot trajectories while still allowing genetic material to be mixed and mutated effectively. The dramatic performance difference between C15 (47.0%) and C1 (76.2%) confirms that proper operator design is more critical than hyperparameter tuning for this problem. Additionally, SUS consistently outperforms roulette selection by maintaining better population diversity, preventing premature convergence, and ensuring more stable evolutionary progress.

3.2.4 Visualization

The visualization component provides real-time feedback on how the Genetic Algorithm (GA) explores the solution space and how the multi-robot paths evolve over time.

Path and Coverage Map Visualization

This visualization shows the environment map, the planned paths of all robots, and the area they have collectively covered. It helps to visually confirm that constraints (like obstacle avoid-

ance and no-collision) are met and provides an intuitive measure of the solution's objective (maximizing coverage).




Figure 3.3: Graphical visualization of the optimization process for the best case study C1, showing robot paths and the area covered in the environment.

Terminal and Metric Visualization

In addition to the graphical output, the terminal provides essential diagnostic information. This output is critical for tracking the results of the GA algorithm.

```
=====
FINAL RESULTS
=====
Visualizing Coverage Map...
=====
Coverage Map (20x20)
=====
 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19
0  .  *  *  *  *  *  *  R2  *  .  .  *  *  *  *  *  *  .  .  .
1  .  *  *  *  *  *  *  *  *  .  *  *  *  *  *  *  *  *  *
2  .  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *
3  .  *  R1  *  *  *  *  *  *  *  *  *  .  *  *  *  *  *  *
4  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *
5  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *
6  .  *  *  .  .  .  .  .  *  *  *  *  *  *  *  *  *  *
7  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *
8  *  *  .  *  *  *  *  *  *  *  .  *  *  *  *  *  *  *
9  *  *  *  *  *  *  *  *  .  *  *  *  *  *  *  *  *  *
10 *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *
11 *  *  R5  *  *  *  *  R6  *  *  *  *  *  *  *  *  *  *
12 .  .  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *
13 .  .  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *
14 *  *  *  .  *  *  *  *  *  *  *  *  *  *  *  *  *  *
15 *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *
16 .  .  .  .  *  *  *  *  *  *  *  *  *  *  *  *  *  *
17 *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *
18 .  .  .  .  *  *  *  *  .  .  R4  *  *  *  *  *  *  *
19 .  .  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *  *
=====
Cells explored: 305/400 (76.2%)
R* = Final robot position, * = Explored, # = Obstacle, . = Unexplored
=====

Cost Breakdown (lower is better):
Coverage count: 305
Connectivity sum: 174.83
Disconnection penalty: 0.00
-----
Coverage cost (α/305): 9.836066
Connectivity cost (β/174.83): 2.859896
Disconnection cost (γ*0.00): 0.000000
-----
Total Cost: 12.695962

Final Best Cost (lower is better): 12.695962
=====
```

Figure 3.4: Terminal visualization and coverage results, displaying key performance metrics.

3.2.5 Conclusions

This study implemented a **Genetic Algorithm** approach for multi-robot area coverage optimization in a 20×20 search-and-rescue grid.

Key Observations

- **Best Performance:** Case C1 achieved the highest coverage (76.2%) using balanced weights and relatively high mutation rate (0.3).
- **Robot Cooperation:** Increasing the number of robots improved total coverage, confirming cooperative exploration benefits.

Parameter Sensitivity

- Coverage performance is highly sensitive to the relative weighting between α and β .
- Increasing robot count consistently improves coverage but introduces additional coordination challenges.

Overall, the study shows that careful tuning of these hyperparameters is essential to maximize multi-robot exploration efficiency using genetic algorithm.

3.3 Ant Colony Optimization Results

3.3.1 Parameters Tuning

The ACO experiments varied pheromone/heuristic weights, evaporation, ant count, and strategy.

Table 3.9: ACO Performance Metrics (AS vs SACO Variants)

Case	Strategy	Robots	Steps	α_τ	β_η	Evap.	Ants	Iter.	Coverage (%)
A1	AS (baseline)	6	125	0.5	1.0	0.40	30	200	93.75
A2	AS (high β_η)	6	125	0.5	2.0	0.35	30	200	88.00
A3	AS (low evap, high α_τ)	6	125	1.0	1.0	0.20	40	200	93.25
A4	SACO (pheromone-only)	6	125	0.7	0.0	0.40	35	200	61.25
A5	SACO (fast evap)	6	125	0.7	0.0	0.60	35	200	58.50

Discussion of Results:

- The AS baseline (A1) achieved the best coverage (93.75%), showing strong balance between pheromone guidance and heuristic bias toward unexplored cells.

- Increasing β_η (A2) reduced coverage to 88.00%, suggesting overly strong heuristic pressure can limit exploration diversity.
- Reducing evaporation and raising α_τ (A3) maintained high coverage (93.25%), indicating slower evaporation can preserve good trails without hurting exploration when α_τ is elevated.
- SACO variants (A4, A5) underperformed markedly (61.25%, 58.50%), highlighting the value of heuristic guidance; faster evaporation further degrades pheromone-only performance.

3.3.2 Artificial Bee Colony Experiments

We ran two ABC configurations headlessly to understand sensitivity to colony size and onlooker ratio. Both reused the shared objective ($\alpha = 3000$, $\beta = 500$, $\gamma = 4.0$) on the 20×20 grid with 6 robots and 100 steps.

Table 3.10: ABC Parameter Sweep (headless runs)

Case	Colony	Cycles	Limit	Onlookers	Window	Attempts	Cost	Coverage (%)
ABC_base	24	200	12	0.5	10	5	14.4231	65.75
ABC_large_colony	32	200	10	0.6	12	6	14.8782	62.25

Discussion:

- The base setting (24 bees, onlookers 0.5) achieved the best cost (14.42) and coverage (65.75%), suggesting diminishing returns from a larger colony in this discrete setting.
- Increasing colony size and onlooker ratio (32 bees, 0.6) slightly worsened cost and coverage, indicating that repair/feasibility pressure dominates over sampling breadth.
- Both runs converged stably; scout resets were crucial to escape stagnation when window perturbations could not find feasible improvements.

3.3.3 Case Study Comparisons (SA vs GA vs ACO vs ABC)

We ran ten case studies with varying team sizes and step counts using the shared cost weights ($\alpha = 3000$, $\beta = 500$, $\gamma = 4.0$) except Case 8 ($\alpha = 2000$, $\beta = 300$). Each algorithm used its standard settings (SA and GA at their tuned defaults; ACO with AS strategy from A1; ABC with the defaults in `case_study_config.py`). Convergence plots are saved under `Figures/case_study-{1..10}-comparison.png`.

Table 3.11: Case Study Results Across SA, GA, ACO, and ABC

Case	Robots	Steps	SA Cost	GA Cost	ACO Cost	ABC Cost	SA Cov (%)	GA Cov (%)	ACO Cov (%)
1	6	125	14.8026	13.3713	10.5770	12.9876	65.50	69.50	70.00
2	4	100	24.6747	21.8845	16.3899	21.2011	39.00	50.00	50.00
3	3	50	63.4858	53.4338	44.3370	49.3614	19.75	21.75	21.75
4	4	100	26.8510	19.7907	16.7808	21.1315	37.25	55.00	55.00
5	4	100	27.1821	21.7872	17.1709	20.3693	39.50	48.25	48.25
6	4	100	27.1199	20.8491	17.2223	21.3523	38.00	50.50	50.50
7	4	100	27.7727	24.4065	16.4758	21.1198	35.75	40.25	40.25
8	4	100	17.0788	14.8692	10.5657	13.1439	40.00	46.50	46.50
9	4	100	32.8435	27.2665	21.3358	20.8085	29.50	37.75	37.75
10	2	100	59.9539	61.1751	39.8666	47.4840	22.00	18.00	18.00

Discussion: Across all ten case studies ACO still delivered the lowest cost and highest coverage in most scenarios, but ABC narrowed the gap: it consistently beat GA on coverage and occasionally on cost (e.g., Case 9) while lagging ACO on coverage. GA remained better than SA on both cost and coverage. Small teams (Cases 3 and 10) remain challenging; ACO leads, ABC sits between ACO and GA/SA. Longer runs (Case 4 with 300 iterations) continue to favor ACO, while ABC maintains mid-pack performance without collapsing coverage. In the convergence figures (Figures 3.5–3.8), the ABC curves trail ACO but often track or overtake GA late in the run, especially when scout resets trigger new feasible regions.

Visual Comparison of Selected Cases

To further illustrate the performance differences, we present the convergence graphs for selected case studies. These graphs plot the cost function value over iterations for SA, GA, and ACO.




Figure 3.5: Convergence comparison for Case 1 (6 Robots, 125 Steps). ACO starts with a superior solution due to heuristic initialization and maintains the lead throughout.




Figure 3.6: Convergence comparison for Case 3 (3 Robots, 50 Steps). ACO demonstrates rapid convergence from a high initial cost to outperform the other algorithms within the first 10 iterations.




Figure 3.7: Convergence comparison for Case 4 (4 Robots, 100 Steps, 300 Iterations). The extended iterations show ACO maintaining its lead and stability.




Figure 3.8: Convergence comparison for Case 10 (2 Robots, 100 Steps). In this challenging scenario with minimal agents, ACO's advantage is particularly pronounced from the start.

The visual results confirm the quantitative findings. ACO consistently achieves the lowest final cost across all scenarios. In cases with larger teams (e.g., Case 1) or specific constraints (e.g., Case 10), ACO benefits significantly from its heuristic initialization, starting with a much lower cost than SA or GA. In other scenarios (e.g., Case 3), ACO may start with a higher cost but demonstrates rapid convergence within the first few iterations to surpass the other algorithms. SA typically shows a sharp initial improvement followed by a plateau, while GA exhibits steadier but more gradual improvement. Ultimately, ACO's combination of pheromone feedback and heuristic guidance provides the most robust performance.

3.3.4 Performance Indices

To rigorously evaluate the optimization performance and consistency of Simulated Annealing (SA), Genetic Algorithm (GA), and Ant Colony Optimization (ACO), we conducted a comprehensive benchmarking study with 20 independent runs of each algorithm. This section presents the comparative analysis of key performance metrics.

Benchmark Methodology

For each algorithm, we performed 20 independent runs using the configuration parameters from Case C1 (6 robots, 125 steps, $\alpha = 3000$, $\beta = 500$). Each run started from a different randomly generated feasible initial solution (for SA and GA) or constructed from scratch (for ACO) to ensure statistical independence. The following metrics were collected:

- **Optimal Solution:** The robot positions that yielded the lowest cost across all runs
- **Optimal Fitness Value:** The minimum cost achieved (lower is better for minimization)
- **Mean Fitness:** Average cost across all 20 runs, indicating typical performance
- **Standard Deviation:** Variance in fitness values, measuring algorithm consistency
- **Computational Time:** Average execution time per run in seconds

Performance Comparison Results

Table 3.11 presents the comparative performance metrics for all three algorithms over 20 independent runs.

Table 3.12: Function Minimization: Performance Indices (20 Runs)

Metric	Simulated Annealing	Genetic Algorithm	Ant Colony Optimization
Optimal Solution	(13, 8), (5, 7), ...	(6, 2), (9, 0), ...	(17, 19), (8, 13) ...
Optimal Fitness Value	13.7049	12.3180	11.4888
Mean Fitness	14.3852	14.3379	11.7522
Standard Deviation	3.6271×10^{-1}	2.1899×10^0	1.1334×10^{-1}
Computational Time	34.867 sec	165.515 sec	130.991 sec

Analysis and Discussion

The benchmark results reveal significant differences in the optimization characteristics of the three algorithms:

- **Solution Quality:** Ant Colony Optimization achieved the best optimal fitness value of 11.4888, outperforming Genetic Algorithm (12.3180) and Simulated Annealing (13.7049).

ACO's ability to construct solutions using both pheromone trails and heuristic information allows it to find superior paths that balance coverage and connectivity more effectively than the other methods.

- **Consistency and Reliability:** ACO demonstrated the highest consistency with the lowest standard deviation (0.1133), compared to SA (0.3627) and GA (2.1899). This indicates that ACO reliably converges to high-quality solutions across different runs, whereas GA exhibits significant variance, likely due to the stochastic nature of its crossover and mutation operators and the random initial population.
- **Computational Efficiency:** Simulated Annealing remains the fastest algorithm (34.867 sec), but its solution quality is the lowest. ACO (130.991 sec) is faster than GA (165.515 sec) while delivering significantly better results. This suggests that ACO offers a better trade-off between solution quality and computational time for this specific multi-robot path planning problem.

3.4 Algorithm Validation

In order to validate our algorithm and experimental results, we benchmarked our performance against the results reported in Kashyap *et al.* (2023) [3]. That study simulated 10 robots operating in a 20 m $\times$ 20 m exploration area over 600 time steps, achieving an average coverage of 76.49 % when starting from random positions.

Our experimental setup uses a comparable area and robot count (in the 10-robot configuration), and by scaling the time step equivalence we aim to ensure that our study is directly comparable.

3.4.1 Time Step Equivalence Calculation

Given the paper's specifications:

- Robot velocity: 0.2 m/s
- Robot size: 0.6m $\times$ 0.6m
- To cover a 3m $\times$ 0.6m grid requires 15 time steps (covering 3 $\times$ 0.6m area)
- To cover the entire 3m $\times$ 3m area requires approximately 75 time steps
- Average time to cover a 1m $\times$ 1m unit cell: 8.33 time steps
- For 600 time steps: average of 72 unit cells covered

Based on this analysis, the original study used 72 and 60 steps for the 10-robot cases to provide comparable experimental conditions.

3.4.2 Results Comparison

To validate our algorithm’s effectiveness, we compare our best-performing configurations with the benchmark results reported by Kashyap *et al.* (2025) [3]. We include the best results from Simulated Annealing (SA), Genetic Algorithm (GA), and Ant Colony Optimization (ACO). For SA and GA, we ran specific 10-robot configurations (C9, C20) to match the benchmark directly. For ACO, we include our best performing case (A1) which used 6 robots but with a longer path length (125 steps), resulting in a comparable total exploration capacity (750 robot-steps vs 720 robot-steps for the benchmark).

Table 3.13: Comparison of Coverage Performance with Benchmark Study

Case	Method	Robots	Steps	α	β	Coverage (%)
C9	SA	10	72	3000	500	67.5
C20	GA	10	72	3000	500	74.5
A1	ACO	6	125	0.5 (α_τ)	1.0 (β_η)	93.75
C21	Benchmark [3]	10	600	–	–	76.49

Discussion of Results:

- **Simulated Annealing (SA):** Achieved 67.5% coverage, approximately 9% lower than the benchmark. This reduction is attributed to the stochastic nature of SA which favored computational efficiency over exhaustive exploration in the limited time steps.
- **Genetic Algorithm (GA):** Achieved 74.5% coverage, coming very close (within 2%) to the benchmark. This demonstrates that the population-based approach effectively explores the solution space, though it requires significantly more computation time than SA.
- **Ant Colony Optimization (ACO):** Achieved the highest coverage of 93.75%, significantly outperforming both the benchmark and the other algorithms. Although this case used fewer robots (6 vs 10), the total exploration capacity was comparable (750 vs 720 robot-steps). The superior performance highlights the effectiveness of the pheromone-based positive feedback loop, which allows agents to cooperatively identify and exploit high-quality paths more efficiently than random search (SA) or evolutionary recombination (GA).

Chapter 4

Conclusion

Chapter 5

Future Work

Appendix

Bibliography

- [1] Zongyuan Shen, Burhanuddin Shirose, Prasanna Sriganesh, Bhaskar Vundurthy, Howie Choset, and Matthew Travers. Multi-cap: A multi-robot connectivity-aware hierarchical coverage path planning algorithm for unknown environments, 2025.
- [2] Jeppe Heini Mikkelsen, Roberto Galeazzi, and Matteo Fumagalli. Optimal multi-robot communication-aware trajectory planning by constraining the fiedler value, 2024.
- [3] Gautam Siddharth Kashyap, Deepkashi Mahajan, Orchid Chetia Phukan, Ankit Kumar, Alexander El Brownlee, and Jiechao Gao. From simulations to reality: enhancing multi-robot exploration for urban search and rescue. *International Journal of Information Technology*, pages 1–12, 2025.
- [4] Suman Som, Marlom Bey, and Pratyay Kuila. Particle swarm optimization for efficient placement of relay nodes in cluster-based wireless sensor networks, 10 2024.
- [5] Xiang Huang, Min Sun, Hang Zhou, and Shuai Liu. A multi-robot coverage path planning algorithm for the environment with multiple land cover types. *IEEE Access*, 8:198101–198117, 2020.